

✓ Clustering Countries for Strategic Aid Allocation

By: Parth Patel

✓ Problem statement:

A humanitarian NGO dedicated to alleviating poverty and providing essential services to underdeveloped nations, has successfully raised \$10 million to support its mission. The organisation now faces the critical task of allocating these funds strategically to maximise impact. The organisation must identify and prioritise the countries in the most dire need of aid, ensuring that limited resources are directed where they can have the greatest positive effect on poverty and community development. The goal is to identify the countries that are most vulnerable and in need of urgent support, enabling the NGO to deploy its resources effectively.

Data description :

Country: Name of the country

Child_mort: Death of children under 5 years of age per 1000 live births

Exports: Exports of goods and services per capita. Given as %age of the GDP per capita

Health: Total health spending per capita. Given as %age of GDP per capita

Imports: Imports of goods and services per capita. Given as %age of the GDP per capita

Income: Net income per person

Inflation: The measurement of the annual growth rate of the Total GDP

Life_expec: The average number of years a new born child would live

Total_fer: The number of children that would be born to each woman


Gdpp: The GDP per capita. Calculated as the Total GDP divided by the total population.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.cluster import AgglomerativeClustering
from sklearn.cluster import DBSCAN
from sklearn.mixture import GaussianMixture
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
from sklearn.metrics import classification_report, accuracy_score
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder, LabelEncoder
from sklearn.feature_selection import RFE
from scipy.stats import uniform
from sklearn.pipeline import make_pipeline
from sklearn.impute import KNNImputer
from sklearn.metrics import roc_curve, auc
from sklearn.neighbors import KNeighborsClassifier
import regex as re
from sklearn.cluster import DBSCAN
from sklearn.cluster import KMeans
from scipy.cluster.hierarchy import dendrogram, linkage
import scipy.stats as stats

import warnings
warnings.filterwarnings('ignore')


df=pd.read_csv("Country-data.csv")
```

```
df.head()
```



	country	child_mort	exports	health	imports	income	inflation	life_expe
0	Afghanistan	90.2	10.0	7.58	44.9	1610	9.44	56
1	Albania	16.6	28.0	6.55	48.6	9930	4.49	76
2	Algeria	27.3	38.4	4.17	31.4	12900	16.10	76
3	Angola	119.0	62.3	2.85	42.9	5900	22.40	60
	Antigua							

```
df.tail()
```



	country	child_mort	exports	health	imports	income	inflation	life_exp
162	Vanuatu	29.2	46.6	5.25	52.7	2950	2.62	65
163	Venezuela	17.1	28.5	4.91	17.6	16500	45.90	75
164	Vietnam	23.3	72.0	6.84	80.2	4490	12.10	75
165	Yemen	56.3	30.0	5.18	34.4	4480	23.60	65
166	Zambia	83.1	37.0	5.89	30.9	3280	14.00	55


```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 167 entries, 0 to 166
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   country         167 non-null    object
1   child_mort       167 non-null    float64
2   exports          167 non-null    float64
3   health           167 non-null    float64
4   imports          167 non-null    float64
5   income           167 non-null    int64
6   inflation        167 non-null    float64
7   life_expec       167 non-null    float64
8   total_fer        167 non-null    float64
9   gdp              167 non-null    int64
dtypes: float64(7), int64(2), object(1)
memory usage: 13.2+ KB
```

```
df.describe(include="all")
```

```
country  child_mort  exports  health  imports  income  inf
count    167      167.000000  167.000000  167.000000  167.000000  167
unique    167           NaN          NaN          NaN          NaN
top  Afghanistan           NaN          NaN          NaN          NaN
freq        1           NaN          NaN          NaN          NaN
mean      NaN      38.270060  41.108976   6.815689  46.890215  17144.688623   7
std      NaN      40.328931  27.412010   2.746837  24.209589  19278.067698  10
min      NaN       2.600000   0.109000   1.810000   0.065900   609.000000  -4
25%      NaN       8.250000  23.800000   4.920000  30.200000  3355.000000   1
50%      NaN      19.300000  35.000000   6.320000  43.300000  9960.000000   5
75%      NaN      62.100000  51.350000   8.600000  58.750000  22800.000000  10
max      NaN     208.000000  200.000000  17.900000  174.000000  125000.000000 104
```

```
sns.heatmap(df.isna())#No Missing Data
```



```
df.duplicated().sum() #No Duplicates
```

 0

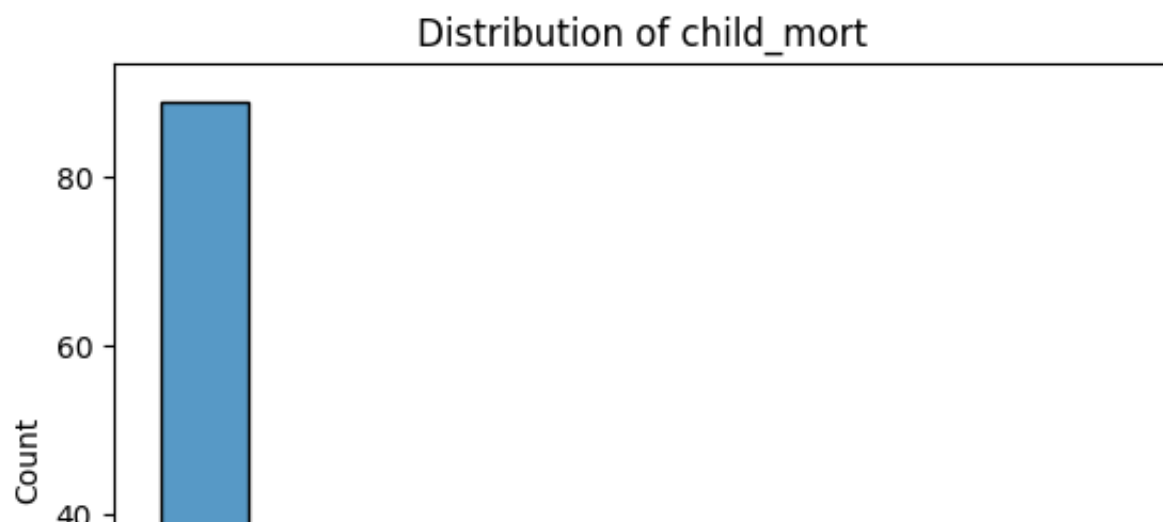
```
df.isna().sum()
```




	0
country	0
child_mort	0
exports	0
health	0
imports	0
income	0
inflation	0
life_expec	0
total_fer	0
gdpp	0


dtype: int64

```
# Univariate Analysis
for col in df.columns:
    if df[col].dtype != 'object':
        plt.figure()
        sns.histplot(df[col])
        plt.title(f"Distribution of {col}")
        plt.show()
```







Distribution of exports




Distribution of health






Distribution of imports




Distribution of income







Distribution of inflation




Distribution of life_expec






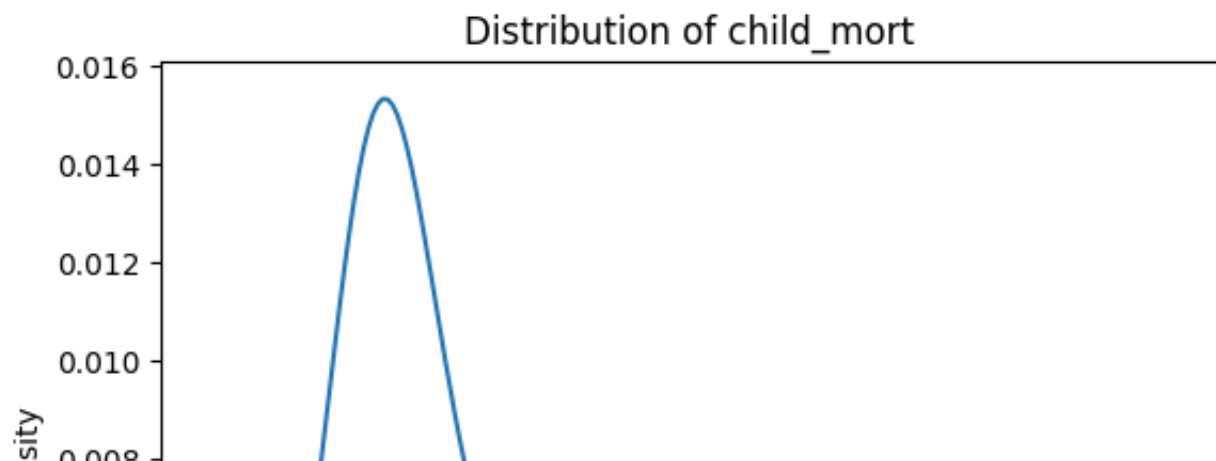
Distribution of total_fer

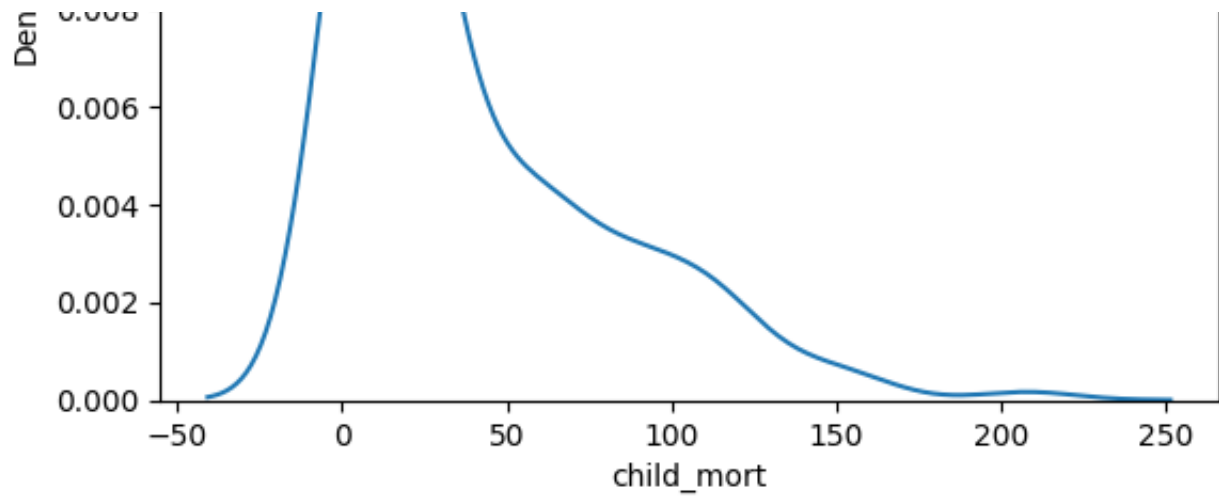


Distribution of gdpp




```
# Univariate Analysis
for col in df.columns:
    if df[col].dtype != 'object':
        plt.figure()
        sns.kdeplot(df[col])
        plt.title(f"Distribution of {col}")
        plt.show()
```







Distribution of exports




Distribution of health







Distribution of imports




Distribution of income







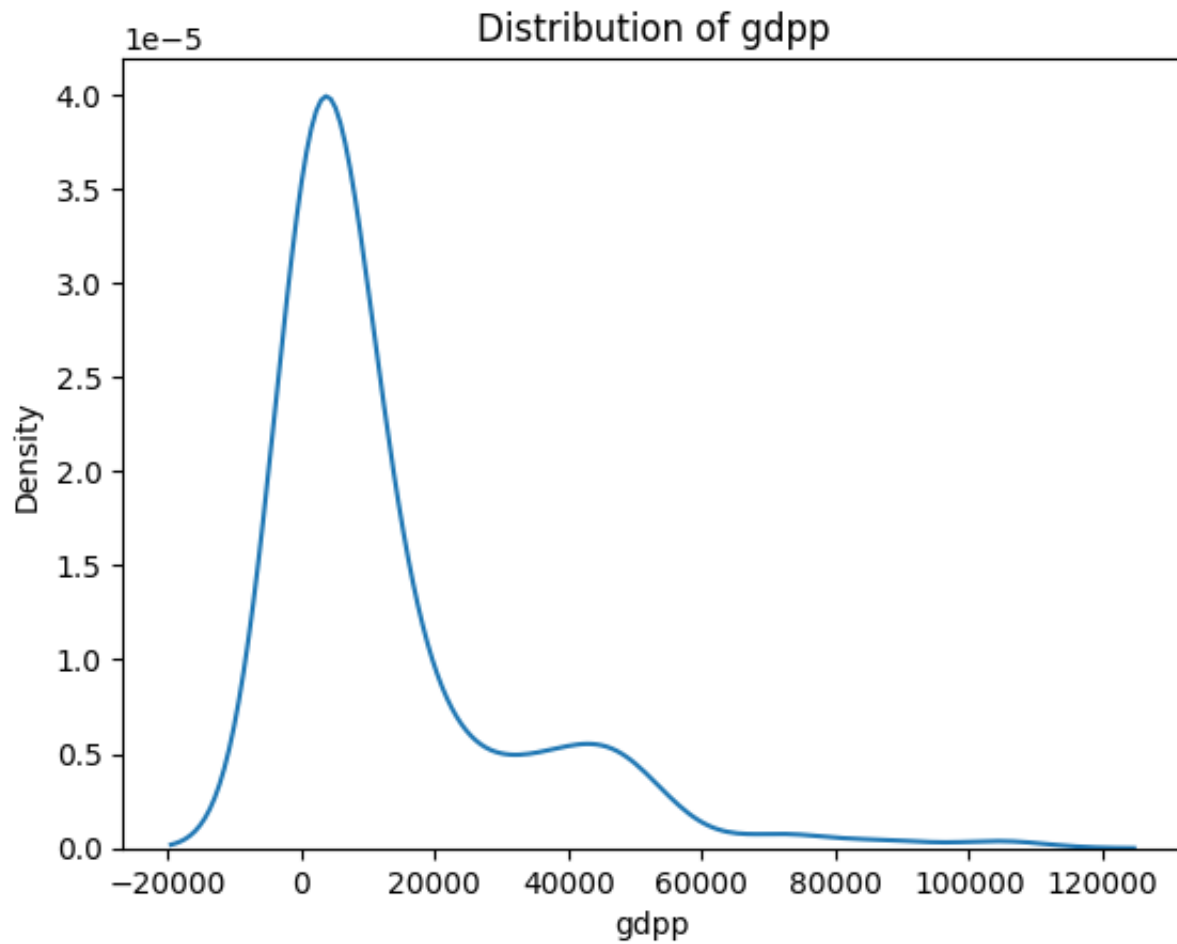
Distribution of inflation



Distribution of life_expec








```
# # Bivariate Analysis  
# sns.pairplot(df)  
# plt.show()
```


#Correlation Heatmap

```
df_corr = df[df.select_dtypes(include=['float', 'int']).columns]
plt.figure(figsize=(12, 8))
sns.heatmap(df_corr.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap of Numerical Columns')
plt.show()
```



```
df['import_export_ratio'] = (df['exports'] / df['imports'])*100

# Outlier Detection using IQR
for col in df.columns:
    if df[col].dtype != 'object':
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - (1.5 * IQR)
        upper_bound = Q3 + (1.5 * IQR)
        outliers = df[(df[col] < lower_bound) | (df[col] > upper_bound)]
        print(f"Outliers in {col}: {outliers.shape[0]}")
```

```
⇒ Outliers in child_mort: 4
   Outliers in exports: 5
   Outliers in health: 2
   Outliers in imports: 4
   Outliers in income: 8
   Outliers in inflation: 5
   Outliers in life_expec: 3
   Outliers in total_fer: 1
   Outliers in gdpp: 25
   Outliers in import_export_ratio: 5
```

```

from scipy.stats.mstats import winsorize

# Apply winsorization to cap outliers
for col in df.columns:
    if df[col].dtype != 'object':
        df[col] = winsorize(df[col], limits=[0.05, 0.05]) # Cap 5% of outliers on bo

# Verify outlier handling (should show reduced or no outliers)
for col in df.columns:
    if df[col].dtype != 'object':
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - (1.5 * IQR)
        upper_bound = Q3 + (1.5 * IQR)
        outliers = df[(df[col] < lower_bound) | (df[col] > upper_bound)]
        print(f"Outliers in {col} after winsorization: {outliers.shape[0]}")#gdp and


```

```

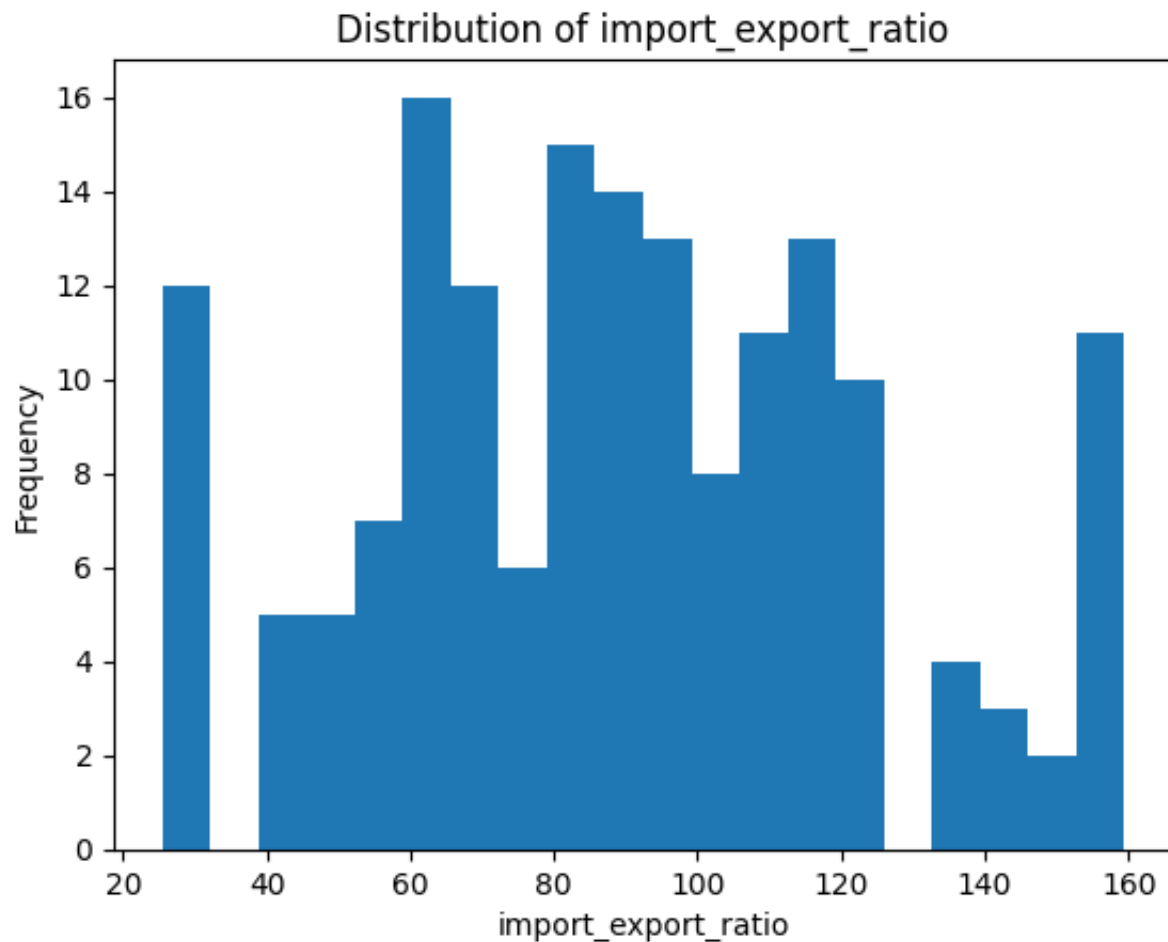
⇒ Outliers in child_mort after winsorization: 0
   Outliers in exports after winsorization: 0
   Outliers in health after winsorization: 0
   Outliers in imports after winsorization: 0
   Outliers in income after winsorization: 0
   Outliers in inflation after winsorization: 0
   Outliers in life_expec after winsorization: 0
   Outliers in total_fer after winsorization: 0
   Outliers in gdpp after winsorization: 25
   Outliers in import_export_ratio after winsorization: 0

```

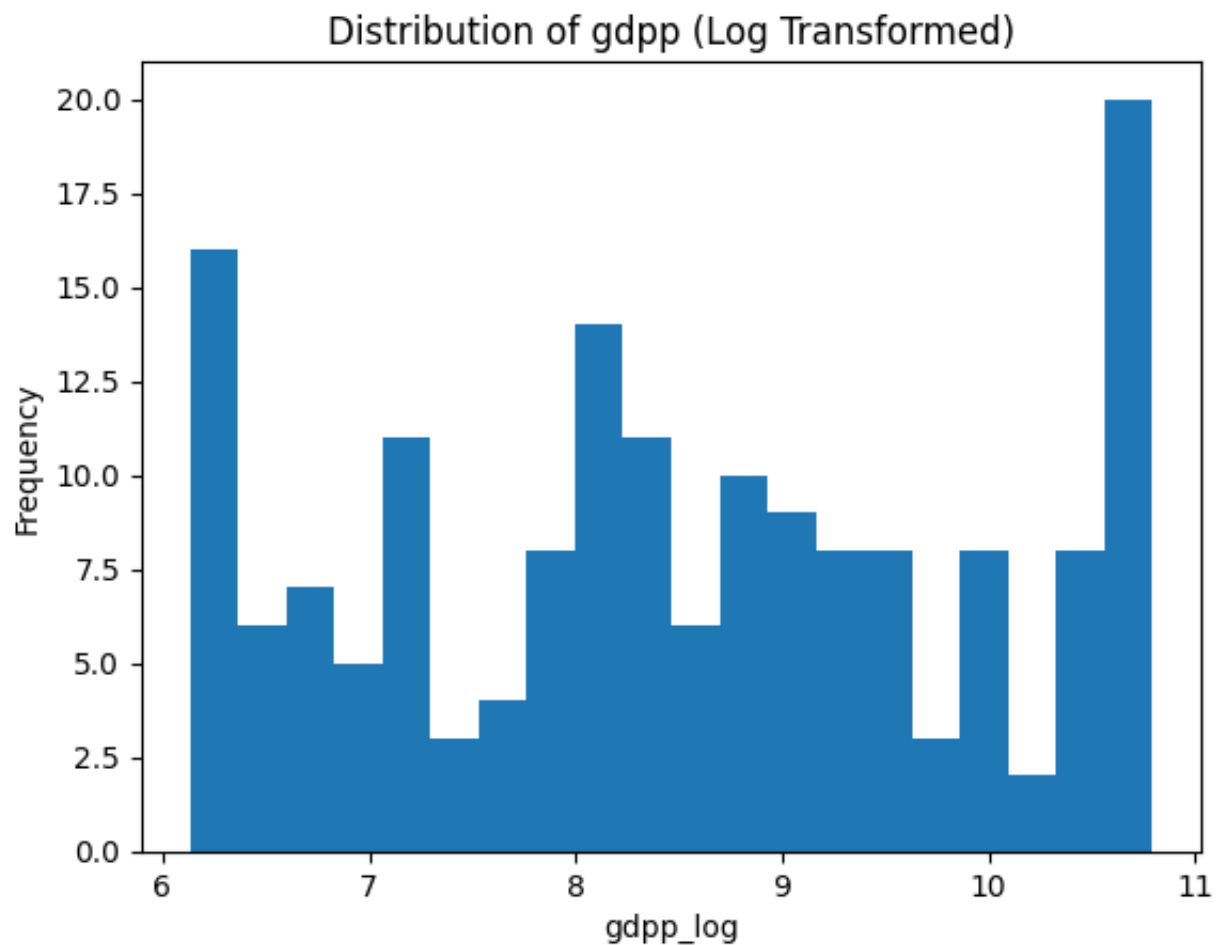
```
# Plot histogram of 'gdpp'  
plt.hist(df['gdpp'], bins=20)  
plt.title('Distribution of gdpp')  
plt.xlabel('gdpp')  
plt.ylabel('Frequency')  
plt.show()
```



```
# Plot histogram of 'import_export_ratio'  
plt.hist(df['import_export_ratio'], bins=20)  
plt.title('Distribution of import_export_ratio')  
plt.xlabel('import_export_ratio')  
plt.ylabel('Frequency')  
plt.show()
```



```
# Applying log transformation to 'gdpp'  
df['gdpp_log'] = np.log1p(df['gdpp'])  
  
# Plot histograms of transformed variables  
plt.hist(df['gdpp_log'], bins=20)  
plt.title('Distribution of gdpp (Log Transformed)')  
plt.xlabel('gdpp_log')  
plt.ylabel('Frequency')  
plt.show()
```



```
# Perform t-test
result = stats.ttest_ind(df['health'], df['life_expec'])

# Print results
print(result)

# Interpret results
alpha = 0.05
if result.pvalue < alpha:
    print("Reject the null hypothesis. Increased health spending (% of GDP) leads
else:
    print("Fail to reject the null hypothesis. There is no significant relationsh

⇒ TtestResult(statistic=-97.99581526818585, pvalue=4.2279500495762e-247, df=332.
Reject the null hypothesis. Increased health spending (% of GDP) leads to high
```

```
# Calculate Pearson correlation coefficient and p-value
corr, p_value = stats.pearsonr(df['health'], df['life_expec'])

# Print results
print(f"Pearson correlation coefficient: {corr}")
print(f"P-value: {p_value}")

# Interpret results
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis. There is a significant correlation between
else:
    print("Fail to reject the null hypothesis. There is no significant correlation

⇒ Pearson correlation coefficient: 0.24947600342499918
P-value: 0.0011488238927149217
Reject the null hypothesis. There is a significant correlation between health
```

```
# Perform t-test
result = stats.ttest_ind(df['total_fer'], df['income'])

# Print results
print(result)

# Interpret results
alpha = 0.05
if result.pvalue < alpha:
    print("Reject the null hypothesis. Countries with higher Total_fertility rate
else:
    print("Fail to reject the null hypothesis. There is no significant relationsh

➡ TtestResult(statistic=-13.709791347305439, pvalue=3.2950919487094535e-34, df=1)
Reject the null hypothesis. Countries with higher Total_fertility rates have
```

```
# Calculate Pearson correlation coefficient and p-value
corr, p_value = stats.pearsonr(df['total_fer'], df['income'])

# Print results
print(f"Pearson correlation coefficient: {corr}")
print(f"P-value: {p_value}")

# Interpret results
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis. There is a significant correlation between
else:
    print("Fail to reject the null hypothesis. There is no significant correlatio

➡ Pearson correlation coefficient: -0.5883441432027093
P-value: 6.239301897978759e-17
Reject the null hypothesis. There is a significant correlation between Total_fer
```



```
# Perform t-test
result = stats.ttest_ind(df['child_mort'], df['income'])

# Print results
print(result)

# Interpret results
alpha = 0.05
if result.pvalue < alpha:
    print("Reject the null hypothesis. Higher income levels are associated with lower child mortality")
else:
    print("Fail to reject the null hypothesis. There is no significant relationship between income and child mortality")

⇒ TtestResult(statistic=-13.680373661937903, pvalue=4.268346375584393e-34, df=30)
Reject the null hypothesis. Higher income levels are associated with lower child mortality
```

```
# Calculate Pearson correlation coefficient and p-value
corr, p_value = stats.pearsonr(df['child_mort'], df['income'])

# Print results
print(f"Pearson correlation coefficient: {corr}")
print(f"P-value: {p_value}")

# Interpret results
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis. There is a significant correlation between income and child mortality")
else:
    print("Fail to reject the null hypothesis. There is no significant correlation between income and child mortality")

⇒ Pearson correlation coefficient: -0.6353145168605605
P-value: 2.925297942217806e-20
Reject the null hypothesis. There is a significant correlation between Child_mortality and income
```

```
# Perform t-test
result = stats.ttest_ind(df['inflation'], df['gdpp'])

# Print results
print(result)

# Interpret results
alpha = 0.05
if result.pvalue < alpha:
    print("Reject the null hypothesis. Higher inflation rates are associated with")
else:
    print("Fail to reject the null hypothesis. There is no significant relationsh")
```

➞ TtestResult(statistic=-10.225772328311542, pvalue=1.6140069109571055e-21, df=)

Reject the null hypothesis. Higher inflation rates are associated with lower (

```
# Calculate Pearson correlation coefficient and p-value
corr, p_value = stats.pearsonr(df['inflation'], df['gdpp'])

# Print results
print(f"Pearson correlation coefficient: {corr}")
print(f"P-value: {p_value}")

# Interpret results
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis. There is a significant correlation between")
else:
    print("Fail to reject the null hypothesis. There is no significant correlation")
```

➞ Pearson correlation coefficient: -0.33293875847535354
P-value: 1.1011891306983982e-05
Reject the null hypothesis. There is a significant correlation between gdpp ar

ML Model

```
df.info()
```

```
↗ <class 'pandas.core.frame.DataFrame'>  
RangeIndex: 167 entries, 0 to 166  
Data columns (total 12 columns):  
#   Column                Non-Null Count  Dtype    
---  -  
0   country                167 non-null   object   
1   child_mort              167 non-null   float64  
2   exports                167 non-null   float64  
3   health                  167 non-null   float64  
4   imports                 167 non-null   float64  
5   income                  167 non-null   int64    
6   inflation               167 non-null   float64  
7   life_expec              167 non-null   float64  
8   total_fer               167 non-null   float64  
9   gdpp                    167 non-null   int64    
10  import_export_ratio     167 non-null   float64  
11  gdpp_log                 167 non-null   float64  
dtypes: float64(9), int64(2), object(1)  
memory usage: 15.8+ KB
```

```
# Define a function to categorize countries into regions
def assign_region(country):
    if country in ['Afghanistan', 'Bangladesh', 'Bhutan', 'India', 'Maldives', 'Nepal']:
        return 'South Asia'
    elif country in ['Brunei', 'Cambodia', 'Indonesia', 'Laos', 'Malaysia', 'Myanmar']:
        return 'Southeast Asia'
    elif country in ['China', 'Hong Kong', 'Japan', 'Macau', 'Mongolia', 'North Korea']:
        return 'East Asia'
    elif country in ['Kazakhstan', 'Kyrgyzstan', 'Tajikistan', 'Turkmenistan', 'Uzbekistan']:
        return 'Central Asia'
    elif country in ['Bahrain', 'Cyprus', 'Egypt', 'Iran', 'Iraq', 'Israel', 'Jordan']:
        return 'Middle East'
    elif country in ['Algeria', 'Angola', 'Benin', 'Botswana', 'Burkina Faso', 'Burundi']:
        return 'Africa'
    elif country in ['Albania', 'Andorra', 'Armenia', 'Austria', 'Azerbaijan', 'Belarus']:
        return 'Europe'
    elif country in ['Antigua and Barbuda', 'Argentina', 'Bahamas', 'Barbados', 'Belize']:
        return 'Americas'
    elif country in ['Australia', 'Fiji', 'Kiribati', 'Marshall Islands', 'Micronesia']:
        return 'Oceania'
    else:
        return 'Other'

# Apply the function to create the 'regions' column
df['regions'] = df['country'].apply(assign_region)
```

```
df.regions.value_counts()
```



	count
regions	
Africa	43
Europe	40
Americas	26
Middle East	15
Southeast Asia	10
Other	9
South Asia	8
Oceania	8
East Asia	4
Central Asia	4

dtype: int64

```
# One Hot Encoding
from sklearn.preprocessing import OneHotEncoder
import pandas as pd

# Ensure the 'regions' column exists
if 'regions' not in df.columns:
    raise KeyError("The column 'regions' does not exist in the DataFrame. Check y

# Ensure no missing values in the 'regions' column
df['regions'] = df['regions'].fillna('Unknown') # Fills missing values if any

# Apply OneHotEncoder
encoder = OneHotEncoder(sparse_output=False, drop='first') # Use sparse_output f
encoded_regions = encoder.fit_transform(df[['regions']])

# Create a DataFrame from the encoded features
encoded_regions_df = pd.DataFrame(
    encoded_regions,
    columns=encoder.get_feature_names_out(['regions'])
```

```
)  
  
# Concatenate the encoded features with the original DataFrame  
df = pd.concat([df, encoded_regions_df], axis=1)  
  
# Drop the original 'regions' column  
df = df.drop('regions', axis=1)  
  
print("One-hot encoding completed successfully!")  
print(df.head()) # Display the updated DataFrame
```

➡ One-hot encoding completed successfully!

	country	child_mort	exports	health	imports	income	\
0	Afghanistan	90.2	12.0	7.58	44.9	1610	
1	Albania	16.6	28.0	6.55	48.6	9930	
2	Algeria	27.3	38.4	4.17	31.4	12900	
3	Angola	116.0	62.3	2.85	42.9	5900	
4	Antigua and Barbuda	10.3	45.5	6.03	58.9	19100	

	inflation	life_expec	total_fer	gdpp	...	gdpp_log	regions_Americas	\
0	9.44	56.2	5.82	553	...	6.317165	0.0	
1	4.49	76.3	1.65	4090	...	8.316545	0.0	
2	16.10	76.5	2.89	4460	...	8.403128	0.0	
3	20.90	60.1	5.87	3530	...	8.169336	0.0	
4	1.44	76.8	2.13	12200	...	9.409273	1.0	

	regions_Central Asia	regions_East Asia	regions_Europe	\
0	0.0	0.0	0.0	
1	0.0	0.0	1.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	regions_Middle East	regions_Oceania	regions_Other	regions_South Asia	\
0	0.0	0.0	0.0	1.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	

	regions_Southeast Asia
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0

[5 rows x 21 columns]

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 167 entries, 0 to 166
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   country                               167 non-null    object
1   child_mort                            167 non-null    float64
2   exports                              167 non-null    float64
3   health                               167 non-null    float64
4   imports                              167 non-null    float64
5   income                               167 non-null    int64
6   inflation                            167 non-null    float64
7   life_expec                           167 non-null    float64
8   total_fer                            167 non-null    float64
9   gdpp                                 167 non-null    int64
10  import_export_ratio                  167 non-null    float64
11  gdpp_log                            167 non-null    float64
12  regions_Americas                    167 non-null    float64
13  regions_Central Asia                 167 non-null    float64
14  regions_East Asia                   167 non-null    float64
15  regions_Europe                      167 non-null    float64
16  regions_Middle East                 167 non-null    float64
17  regions_Oceania                     167 non-null    float64
18  regions_Other                       167 non-null    float64
19  regions_South Asia                  167 non-null    float64
20  regions_Southeast Asia              167 non-null    float64
dtypes: float64(18), int64(2), object(1)
memory usage: 27.5+ KB
```

```
# Select numerical columns for scaling
numerical_cols = df.select_dtypes(include=['float', 'int']).columns

# Initialize the scaler
scaler = StandardScaler() # or MinMaxScaler() depending on your preference


# Fit and transform the numerical columns
df[numerical_cols] = scaler.fit_transform(df[numerical_cols])
```

```
import matplotlib.pyplot as plt
# Assuming 'df' is your DataFrame with relevant features for clustering

# Select numerical features for clustering
X = df[['child_mort', 'exports', 'health', 'imports', 'income', 'inflation',
```



```
'life_expec', 'total_fer', 'import_export_ratio', 'gdpp_log',  
'regions_Americas', 'regions_Central Asia', 'regions_East Asia',  
'regions_Europe', 'regions_Middle East', 'regions_Oceania',  
'regions_Other', 'regions_South Asia', 'regions_Southeast Asia']]  
  
# Initialize list to store WCSS (Within-Cluster Sum of Squares)  
wcss = []  
  
# Try different values of k (number of clusters)  
for i in range(1, 11):  
    kmeans = KMeans(n_clusters=i, random_state=42)  
    kmeans.fit(X)  
    wcss.append(kmeans.inertia_)  
  
# Plot the Elbow Method graph  
plt.plot(range(1, 11), wcss, marker='o')  
plt.title('Elbow Method for Optimal k')  
plt.xlabel('Number of Clusters (k)')  
plt.ylabel('WCSS')  
plt.show()
```



```
# Based on the Elbow method, choose the optimal k (let's assume k=3)
optimal_k = 3

# Initialize KMeans with the optimal k
kmeans = KMeans(n_clusters=optimal_k, random_state=42)

# Fit the model to the data
kmeans.fit(X)


# Get cluster labels for each data point
df['Cluster'] = kmeans.labels_

# Print the cluster centers
print(kmeans.cluster_centers_)

# Visualize the clusters (example using two features)
plt.scatter(df['imports'], df['exports'], c=df['Cluster'], cmap='viridis')
plt.title('K-Means Clustering')
plt.xlabel('Imports')
```

```
plt.ylabel('Exports')
plt.show()
```

```
[[ 9.26227057e-01 -5.03722974e-01 -2.74155474e-01 -1.87725906e-04
 -8.22984564e-01  2.47450284e-01 -9.71064453e-01  9.03330152e-01
 -5.94518621e-01 -9.52721222e-01 -2.64748342e-01  3.85784998e-02
 -5.90367951e-02 -5.61213533e-01 -3.14140431e-01  1.95014794e-01
  2.24022598e-01  2.64902070e-01 -7.53364873e-04]
 [-7.28442486e-01  1.52517322e-01  6.02429274e-01  1.91612419e-02
  5.25632345e-01 -4.93780426e-01  7.45235904e-01 -7.72863163e-01
  1.65374293e-01  6.95383809e-01  3.98034588e-01 -1.56652090e-01
  1.23643257e-01  7.77681610e-01 -3.14140431e-01 -9.05246477e-02
 -1.12135693e-01 -2.24308862e-01 -2.52377233e-01]
 [-3.68874627e-01  7.69107557e-01 -7.93387749e-01 -4.42903100e-02
  6.11523389e-01  5.99515359e-01  4.29826835e-01 -2.14089958e-01
  9.41884901e-01  5.05181843e-01 -3.37476075e-01  2.79362894e-01
 -1.56652090e-01 -5.61213533e-01  1.43457464e+00 -2.24308862e-01
 -2.38667185e-01 -6.82272787e-02  5.90562724e-01]]
```



```
# Create a dendrogram to visualize the hierarchical structure
import scipy.cluster.hierarchy as sch
```

```

import matplotlib.pyplot as plt
from sklearn.cluster import AgglomerativeClustering


# Dendrogram to determine the number of clusters
dendrogram = sch.dendrogram(sch.linkage(X, method='ward'))
plt.title('Dendrogram')
plt.xlabel('Data Points')
plt.ylabel('Euclidean Distances')
plt.show()


# Perform Agglomerative Clustering (choose the number of clusters based on the dendrogram)
# Let's assume we want 3 clusters
hc = AgglomerativeClustering(n_clusters=3, metric='euclidean', linkage='ward') #
y_hc = hc.fit_predict(X)

# Add cluster labels to your DataFrame
df['Cluster_HC'] = y_hc

# Visualize the clusters (example using two features)
plt.scatter(df['imports'], df['exports'], c=df['Cluster_HC'], cmap='viridis')
plt.title('Hierarchical Clustering')
plt.xlabel('Imports')
plt.ylabel('Exports')
plt.show()

```






```
# Initialize DBSCAN
dbscan = DBSCAN(eps=0.5, min_samples=5) # Adjust eps and min_samples as needed

# Fit the model to the data
y_dbscan = dbscan.fit_predict(X)

# Add cluster labels to your DataFrame
df['Cluster_DBSCAN'] = y_dbscan

# Visualize the clusters (example using two features)
plt.scatter(df['imports'], df['exports'], c=df['Cluster_DBSCAN'], cmap='viridis')
plt.title('DBSCAN Clustering')
plt.xlabel('Imports')
plt.ylabel('Exports')
plt.show()
```



```

from sklearn.metrics import silhouette_score

# Calculate Silhouette Coefficient for KMeans
silhouette_kmeans = silhouette_score(X, df['Cluster'])
print("Silhouette Coefficient for KMeans:", silhouette_kmeans)


# Calculate Silhouette Coefficient for Agglomerative Clustering
silhouette_hc = silhouette_score(X, df['Cluster_HC'])
print("Silhouette Coefficient for Hierarchical Clustering:", silhouette_hc)

# Calculate Silhouette Coefficient for DBSCAN (excluding noise points)
core_samples_mask = np.zeros_like(y_dbscan, dtype=bool)
core_samples_mask[dbscan.core_sample_indices_] = True
labels_dbscan = y_dbscan[core_samples_mask]
X_dbscan = X[core_samples_mask]
if len(set(labels_dbscan)) > 1: # Check if there are at least two clusters (excl
    silhouette_dbscan = silhouette_score(X_dbscan, labels_dbscan)
    print("Silhouette Coefficient for DBSCAN:", silhouette_dbscan)
else:
    print("Silhouette Coefficient for DBSCAN cannot be calculated as there are le

⇒ Silhouette Coefficient for KMeans: 0.1998189829076175
   Silhouette Coefficient for Hierarchical Clustering: 0.18462409237755378
   Silhouette Coefficient for DBSCAN cannot be calculated as there are less than

```

```
# Visualize the clusters (example using two features)
plt.scatter(df['imports'], df['exports'], c=df['Cluster'], cmap='viridis')
plt.title('Clustering Results')
plt.xlabel('Imports')
plt.ylabel('Exports')
plt.show()
```



```
from sklearn.decomposition import PCA

# Select numerical features for PCA
X = df.select_dtypes(include=['float', 'int'])

# Initialize PCA with 2 components
pca = PCA(n_components=2)


# Fit and transform the data
X_pca = pca.fit_transform(X)
```



```
# Create a DataFrame from the PCA results
pca_df = pd.DataFrame(data=X_pca, columns=['PC1', 'PC2'])

# Add cluster labels to the PCA DataFrame
pca_df['Cluster'] = df['Cluster'] # Replace 'Cluster' with the actual column name

# Visualize the clusters using PCA
plt.scatter(pca_df['PC1'], pca_df['PC2'], c=pca_df['Cluster'], cmap='viridis')
plt.title('Clustering Results Visualized with PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```



```
# Get the cluster centers
cluster_centers = kmeans.cluster_centers_

# Calculate the absolute values of the cluster centers
abs_centers = np.abs(cluster_centers)
```

```
# Sort features within each cluster based on their absolute values in the cluster
ordered_features = []
for i in range(3): # Assuming optimal_k is the number of clusters
    ordered_features.append(np.argsort(abs_centers[i])[:, :-1])

# Print the most important features for each cluster
for i in range(3):
    print(f"Cluster {i}:")
    for feature_idx in ordered_features[i]:
        print(f"  - {X.columns[feature_idx]}")
```

➡ Cluster 0:

- life_expec
- import_export_ratio
- child_mort
- total_fer
- income
- gdpp
- regions_East Asia
- exports
- regions_Europe
- health
- regions_Other
- gdpp_log
- inflation
- regions_Oceania
- regions_Middle East
- regions_Central Asia
- regions_Americas
- regions_South Asia
- imports

Cluster 1:

- regions_East Asia
- total_fer
- life_expec
- child_mort
- import_export_ratio
- health
- income
- inflation
- gdpp_log
- regions_Europe
- regions_South Asia
- regions_Other
- gdpp
- regions_Americas

- exports
- regions_Central Asia
- regions_Oceania
- regions_Middle East
- imports

Cluster 2:

- regions_Europe
- gdpp
- health
- exports
- income
- inflation
- regions_South Asia
- regions_East Asia
- import_export_ratio
- life_expec
- child_mort
- gdpp_log
- regions_Americas
- regions_Oceania
- regions_Middle East
- total_fer
- regions_Central Asia
- regions_Other

```

# Get the features of the given country
country_name = input("Enter a country") # Replace with the actual country name
country_data = df[df['country'] == country_name][['child_mort', 'exports', 'health_
    'life_expec', 'total_fer', 'import_export_ratio', 'gdpp_log',
    'regions_Americas', 'regions_Central Asia', 'regions_East Asia',
    'regions_Europe', 'regions_Middle East', 'regions_Oceania',
    'regions_Other', 'regions_South Asia', 'regions_Southeast Asia']].values

# Predict the cluster for the given country
country_cluster = kmeans.predict(country_data)[0]

# Find countries in the same cluster
similar_countries = df[df['Cluster'] == country_cluster]['country'].tolist()

# Remove the given country from the list of similar countries
similar_countries.remove(country_name)

print(f"Countries similar to {country_name} based on KMeans clustering: {similar_

```

↔ Enter a country India
 Countries similar to India based on KMeans clustering: ['Afghanistan', 'Angola',

```

import pickle
import os
#File Path
file_path = 'Clustering Countries for Strategic Aid Allocation.pkl'

# Save the KMeans model to the pickle file
with open(file_path, 'wb') as file:
    pickle.dump(kmeans, file)

```

Recommendations

1. Invest in Healthcare:

- Focus on reducing high child mortality rates by funding immunization programs, enhancing maternal healthcare, and improving basic health infrastructure. These initiatives are critical for saving lives and fostering a healthier future for vulnerable populations.

2. Collaborate to Strengthen Local Health Systems:

- Partner with local organizations to enhance healthcare facilities, supply medical resources, and train healthcare workers in regions with low life expectancy. A community-driven approach ensures sustainable improvements in health outcomes.

3. Prioritize Long-Term Development Projects:

- Allocate funds toward clean water access, education, and renewable energy projects to address the root causes of poverty and foster long-term economic and social stability.

4. Support Education and Skill Development:

- Promote early childhood education, health awareness campaigns, and vocational training programs in underserved areas. By providing scholarships and enhancing teacher training, these initiatives can empower communities and drive sustainable change.

5. Enhance Nutrition and Food Security:

- Address high child mortality rates through sustainable agriculture programs, nutritional education, and food supplements. Implementing school meal initiatives can improve both health and educational outcomes, laying a strong foundation for future progress.

